

You should run your program with  
`mpirun -n N` (executable name)  
 where the number of processors  $N$  is a square, say  $N = p^2$ .

For an  $n \times n$  matrix  $M$ , we assume  $M$  to be logically divided into  $p^2$  square submatrices, each of size  $(n/p) \times (n/p)$ . We denote the  $(i, j)$ -th processor by  $P_{i,j}$  and  $(i, j)$ -th submatrix of  $M$  by  $M_{i,j}$  where  $0 \leq i, j < p$ . Initially  $P_{i,j}$  should generate  $M_{i,j}$ , so that at the beginning of the given algorithm outline we assume that  $M_{i,j}$  is already available to  $P_{i,j}$

## Cannon's Matrix Multiplication (with Non-Blocking Communication Primitives)

Notation: For an integer  $x$ ,  $\bar{x}$  is the integer s.t.  $0 \leq \bar{x} < p$  and  $x \equiv \bar{x} \pmod{p}$

### Algorithm outline

Goal: Given  $n \times n$  matrices  $A$  and  $B$  compute  $C = AB$ .

Initially  $P_{i,j}$  generates  $A_{i,j}$  and  $B_{i,j}$

Perform cyclic shift on  $i$ -th row of  $A$  by  $i$  places so that  $P_{i,j}$  contains  $A_{i,\bar{i+j}}$

Perform cyclic shift on  $j$ -th column of  $B$  by  $j$  places so that  $P_{i,j}$  contains  $B_{\bar{i+j},j}$

( Notice that at the beginning of  $k$ -th iteration,  $P_{i,j}$  contains  $A_{i,\bar{i+j+k}}$  and  $B_{\bar{i+j+k},j}$  )

For  $k = 0, \dots, p-1$  do the following in  $P_{i,j}$

    If  $k < p-1$

        Initiate cyclic shift on all rows of  $A$  by 1 place

        Initiate cyclic shift on all columns of  $B$  by 1 place

    If  $k = 0$ , compute  $C_{i,j} = A_{i,\bar{i+j}} B_{\bar{i+j},j}$

    If  $k > 0$ , compute  $C_{i,j} = C_{i,j} + A_{i,\bar{i+j+k}} B_{\bar{i+j+k},j}$

    Wait for cyclic shifts to complete so that,  $P_{i,j}$  contains  $A_{i,\bar{i+j+k+1}}$  and  $B_{\bar{i+j+k+1},j}$

( Notice that at the end of  $p-1$ -th iteration,  $P_{i,j}$  contains  $C_{i,j}$  is indeed the submatrix corresponding to  $C = AB$  )

### Note

Perform the cyclic shift inside the loop using non-blocking primitives. Also, in round  $k$ ,  $P_{i,j}$  performs computation with  $A_{i,\bar{i+j+k}}$  and  $B_{\bar{i+j+k},j}$  while receiving  $A_{i,\bar{i+j+k+1}}$  and  $B_{\bar{i+j+k+1},j}$  from  $P_{i,\bar{j+1}}$  and  $P_{\bar{i+1},j}$  respectively. Hence you need to use **different buffers** to hold the data that you are computing with and the data that you are sending/receiving.